

Indexes – example

```

use Example_Lab1
go

create table Tea(
Tid int primary key identity,
TName varchar(50),
Price int)

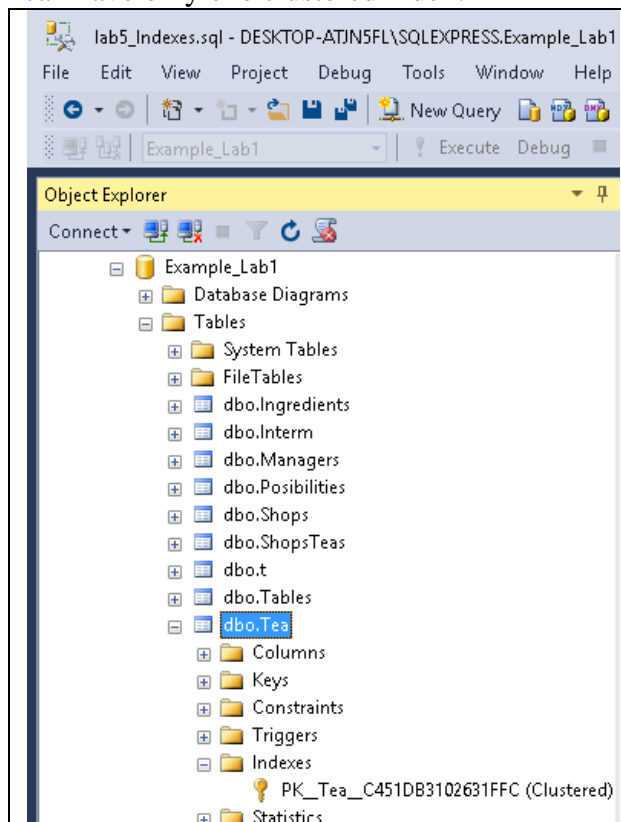
insert into Tea values ('Mint', 10),
('Ginger', 12), ('Fruits', 9), ('Rose', 8)

select * from Tea

```

| | Tid | TName | Price |
|---|-----|--------|-------|
| 1 | 1 | Mint | 10 |
| 2 | 2 | Ginger | 12 |
| 3 | 3 | Fruits | 9 |
| 4 | 4 | Rose | 8 |

Automatically a clustered index is created on the primary key (when this one is created). On a table one can have only one clustered index.

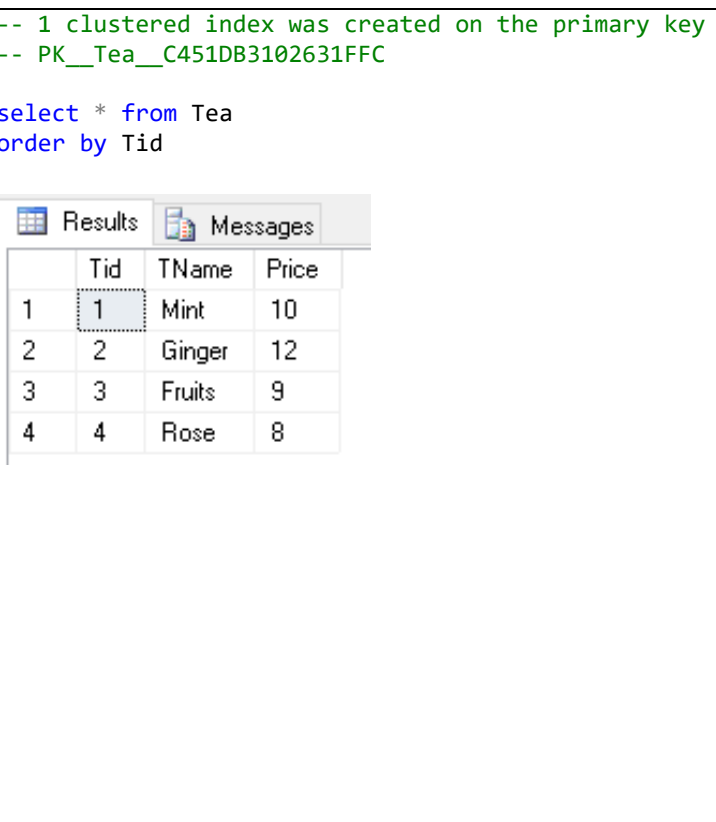


```

-- 1 clustered index was created on the primary key
-- PK__Tea__C451DB3102631FFC

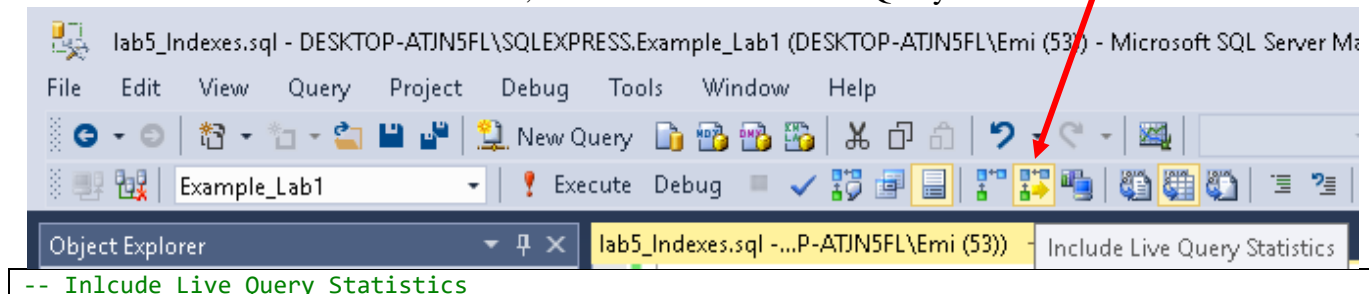
select * from Tea
order by Tid

```



| | Tid | TName | Price |
|---|-----|--------|-------|
| 1 | 1 | Mint | 10 |
| 2 | 2 | Ginger | 12 |
| 3 | 3 | Fruits | 9 |
| 4 | 4 | Rose | 8 |

To check the indexes and how are used, one can use Include Live Query Statistics.



```

-- Include Live Query Statistics

```

```
select * from Tea
order by Tid
```

```
-- Include Live Query Statistics
select * from Tea
order by Tid
```

100 %

Results Messages Live Query Statistics

Estimated query progress: 100% Query 1: Query cost (relative to the batch): 100% select * from Tea order by Tid

SELECT

Clustered Index Scan (Clustered)

[Tea].[PK_Tea_C451DB3102631FFC]

4 of 4 (100%)

Clustered Index Scan (Clustered)

Scanning a clustered index, entirely or only a range.

Estimated operator progress: 100%

| | |
|--------------------------------|----------------------|
| Physical Operation | Clustered Index Scan |
| Logical Operation | Clustered Index Scan |
| Actual Execution Mode | Row |
| Estimated Execution Mode | Row |
| Storage | RowStore |
| Number of Rows Read | 4 |
| Actual Number of Rows | 4 |
| Actual Number of Batches | 0 |
| Estimated I/O Cost | 0.003125 |
| Estimated Operator Cost | 0.0032864 (100%) |
| Estimated Subtree Cost | 0.0032864 |
| Estimated CPU Cost | 0.0001614 |
| Estimated Number of Executions | 1 |
| Number of Executions | 1 |
| Estimated Number of Rows | 4 |
| Estimated Row Size | 44 B |
| Actual Rebinds | 0 |
| Actual Rewinds | 0 |
| Ordered | True |
| Node ID | 0 |

Object


[Example_Lab1].[dbo].[Tea].[PK_Tea_C451DB3102631FFC]

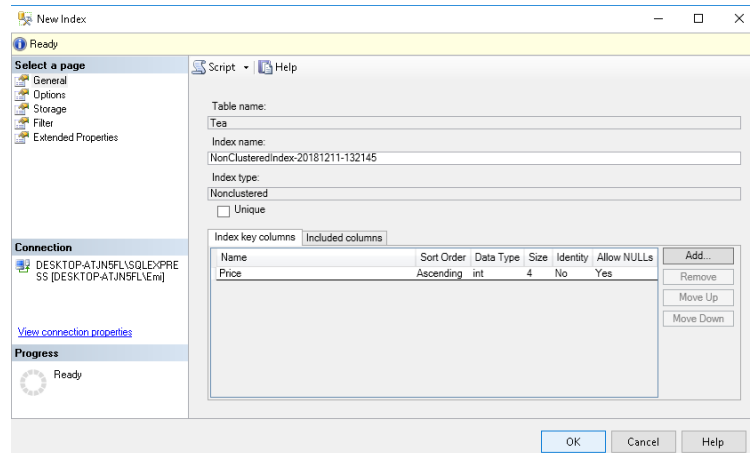
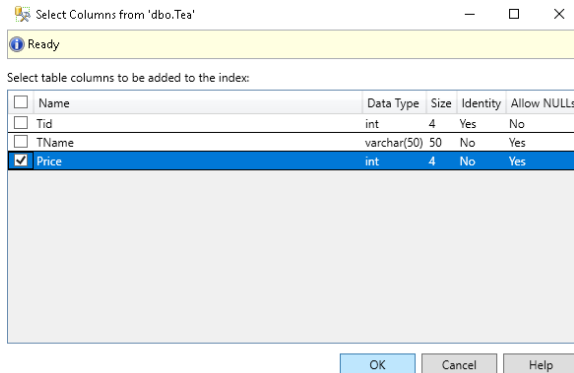
Output List

[Example_Lab1].[dbo].[Tea].Tid, [Example_Lab1].[dbo].[Tea].TName, [Example_Lab1].[dbo].[Tea].Price

Create Non-Clustered Indexes by Design View

- in the table tabs – choose Indexes – right click – new Index – Non-Clustered Index – Name (if one wants to change it) – Add (choose the field(s) for the non-clustered index) – ok - ok





```
-- create non-clustered index - by Design View
-- NonClusteredIndex-20181211-132145
```

Create Non-Clustered Indexes by Code

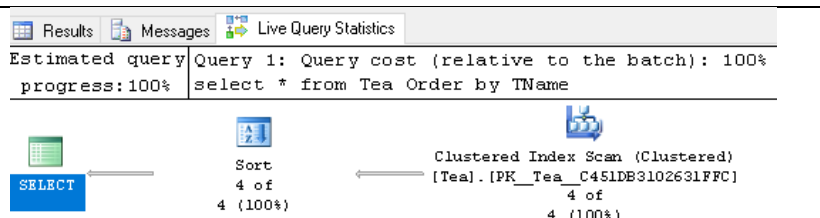
```
-- all the indexes from the
current database
select * from sys.indexes
```

| object_id | name | index_id | type | type_desc | is_unique | data_space_id | ignore_dup_key | is_primary_key | is_unique_constraint | fill_factor | is_padded |
|-----------|------|----------|------|-----------|--------------|---------------|----------------|----------------|----------------------|-------------|-----------|
| 1 | 3 | clst | 1 | 1 | CLUSTERED | 1 | 1 | 0 | 0 | 0 | 0 |
| 2 | 5 | clst | 1 | 1 | CLUSTERED | 1 | 1 | 0 | 0 | 0 | 0 |
| 3 | 6 | clst | 1 | 1 | CLUSTERED | 1 | 1 | 0 | 0 | 0 | 0 |
| 4 | 7 | clst | 1 | 1 | CLUSTERED | 1 | 1 | 0 | 0 | 0 | 0 |
| 5 | 7 | nc | 2 | 2 | NONCLUSTERED | 1 | 1 | 0 | 0 | 0 | 0 |
| 6 | 8 | NULL | 0 | 0 | HEAP | 0 | 1 | 0 | 0 | 0 | 0 |
| 7 | 9 | clst | 1 | 1 | CLUSTERED | 1 | 1 | 0 | 0 | 0 | 0 |
| 8 | 17 | cl | 1 | 1 | CLUSTERED | 1 | 1 | 0 | 0 | 0 | 0 |
| 9 | 17 | nc | 2 | 2 | NONCLUSTERED | 1 | 1 | 0 | 0 | 0 | 0 |
| 10 | 17 | nc2 | 3 | 2 | NONCLUSTERED | 1 | 1 | 0 | 0 | 0 | 0 |
| 11 | 18 | cl | 1 | 1 | CLUSTERED | 1 | 1 | 0 | 0 | 0 | 0 |

```
select name from sys.indexes
```

| name |
|------|
| 1 |
| 2 |
| 3 |
| 4 |
| 5 |
| 6 |
| 7 |
| 8 |
| 9 |
| 10 |
| 11 |

```
select * from Tea
Order by TName
-- only the clustered index is used
```



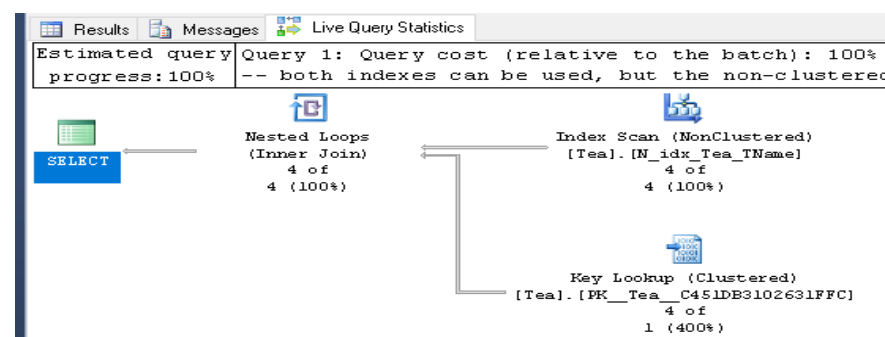
| Clustered Index Scan (Clustered) | |
|--|----------------------|
| Scanning a clustered index, entirely or only a range. | |
| Estimated operator progress: 100% | |
| Physical Operation | Clustered Index Scan |
| Logical Operation | Clustered Index Scan |
| Actual Execution Mode | Row |
| Estimated Execution Mode | Row |
| Storage | RowStore |
| Number of Rows Read | 4 |
| Actual Number of Rows | 4 |
| Actual Number of Batches | 0 |
| Estimated I/O Cost | 0.003125 |
| Estimated Operator Cost | 0.0032864 (22%) |
| Estimated CPU Cost | 0.0001614 |
| Estimated Subtree Cost | 0.0032864 |
| Estimated Number of Executions | 1 |
| Number of Executions | 1 |
| Estimated Number of Rows | 4 |
| Estimated Row Size | 44 B |
| Actual Rebids | 0 |
| Actual Rewinds | 0 |
| Ordered | False |
| Node ID | 1 |
| Object | |
| [Example_Lab1].[dbo].[Tea].[PK_Tea_C451DB3102631FFC] | |
| Output List | |
| [Example_Lab1].[dbo].[Tea].Tid, [Example_Lab1].[dbo].[Tea].TName, [Example_Lab1].[dbo].[Tea].Price | |

```
--create index non-clustered on the TName field
IF EXISTS (SELECT name FROM sys.indexes WHERE name = N'N_idx_Tea_TName')
    DROP INDEX N_idx_Tea_TName ON Tea;
GO

CREATE NONCLUSTERED INDEX N_idx_Tea_TName ON Tea(TName);
GO

-- both indexes can be used, but the non-clustered one is more efficient
select * from Tea
Order by TName
```

| Index Scan (NonClustered) | |
|--|-----------------|
| Scan a nonclustered index, entirely or only a range. | |
| Estimated operator progress: 100% | |
| Physical Operation | Index Scan |
| Logical Operation | Index Scan |
| Estimated Execution Mode | Row |
| Storage | RowStore |
| Actual Number of Rows | 4 |
| Estimated Operator Cost | 0.0032864 (47%) |
| Estimated I/O Cost | 0.003125 |
| Estimated CPU Cost | 0.0001614 |
| Estimated Subtree Cost | 0.0032864 |
| Number of Executions | 1 |
| Estimated Number of Executions | 1 |
| Estimated Number of Rows | 4 |
| Estimated Row Size | 40 B |
| Ordered | True |
| Node ID | 1 |
| Object | |
| [Example_Lab1].[dbo].[Tea].[N_idx_Tea_TName] | |
| Output List | |
| [Example_Lab1].[dbo].[Tea].Tid, [Example_Lab1].[dbo].[Tea].TName | |



| Key Lookup (Clustered) | |
|---|-----------------|
| Uses a supplied clustering key to lookup on a table that has a clustered index. | |
| Estimated operator progress: 100% | |
| Physical Operation | Key Lookup |
| Logical Operation | Key Lookup |
| Estimated Execution Mode | Row |
| Storage | RowStore |
| Actual Number of Rows | 4 |
| Estimated Operator Cost | 0.0037574 (53%) |
| Estimated I/O Cost | 0.003125 |
| Estimated CPU Cost | 0.0001581 |
| Estimated Subtree Cost | 0.0037574 |
| Number of Executions | 4 |
| Estimated Number of Executions | 4 |
| Estimated Number of Rows | 1 |
| Estimated Row Size | 11 B |
| Ordered | True |
| Node ID | 3 |
| Object | |
| [Example_Lab1].[dbo].[Tea]. | |
| [PK_Tea_C451DB3102631FFC] | |
| Output List | |
| [Example_Lab1].[dbo].[Tea].Price | |
| Seek Predicates | |
| Seek Keys[1]: Prefix: [Example_Lab1].[dbo].[Tea].Tid = | |
| Scalar Operator([Example_Lab1].[dbo].[Tea].[Tid]) | |

dbo.Tea

- Columns
- Keys
- Constraints
- Triggers
- Indexes
 - N_idx_Tea_TName (Non-Unique, Non-Clustered)
 - NonClusteredIndex-20181211-132145 (Non-Unique, Non-Clustered)
 - PK_Tea_C451DB3102631FFC (Clustered)
- Statistics

The Non-clustered index should be created on the fields involved in ORDER BY clauses, WHERE clause, JOIN clauses, to increase the efficiency and decrease the execution time.

Create view

```
-- create view
create view vTea
as
    select * from Tea
    where TName LIKE 'a%'
go

-- execute
select * from vTea
order by TName
```

